

EC3357:Machine Learning

Lecture 1: Introduction to Machine Learning

What is Learning?

- Herbert Simon: “Learning is any process by which a system improves performance from experience.”
- “A computer program is said to **learn** from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E.” – Tom Mitchell
- From the definition by Tom Mitchell (1998):
 - Machine Learning is the study of algorithms that
 - improve their performance P
 - at some task T
 - with experience E.
 - A well-defined learning task is given by $\langle P, T, E \rangle$.

Traditional Programming



Machine Learning



How can we solve a specific problem?

- As computer scientists we write a program that encodes a set of rules that are useful to solve the problem
- In many cases is very difficult to specify those rules, e.g., given a picture determine whether there is a cat in the image.



Learning

- Learning systems are not directly programmed to solve a problem, instead develop own program based on:
 - Examples of how they should behave
 - From trial-and-error experience trying to solve the problem
- Learning simply means incorporating information from the training(input-output pairs) examples into the system

A classic example of a task that requires machine learning

- It is very hard to say what makes a 2



Why use learning?

- It is very hard to write programs that solve problems like recognizing a handwritten digit
 - What distinguishes a 2 from a 7?
 - How does our brain do it?
- Instead of writing a program by hand, we collect examples that specify the correct output for a given input
- A machine learning algorithm then takes these examples and produces a program that does the job

When Do We Use Machine Learning

- ML is used when:
 - Human expertise does not exist (navigating on Mars)
 - Humans can't explain their expertise (speech recognition)
 - Models must be customized (personalized medicine)
 - Models are based on huge amounts of data (genomics)



- Learning isn't always useful:
 - There is no need to “learn” to calculate payroll

More examples of tasks that are best solved by using a learning algorithm

- Recognizing patterns:
 - Facial identities or facial expressions
 - Handwritten or spoken words
 - Medical images
- Generating patterns:
 - Generating images or motion sequences
- Recognizing anomalies:
 - Unusual credit card transactions
 - Unusual patterns of sensor readings in a nuclear power plant
- Prediction:
 - Future stock prices or currency exchange rates

Sample Applications

- Web search
 - Computational biology
 - Finance
 - E-commerce
 - Space exploration
 - Robotics
 - Information extraction
 - Social networks
 - Debugging software
 - [Your favorite area]

Samuel's Checkers-Player

- “Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.” - Arthur Samuel (1959)



Defining the Learning Task

Improve on task T, with respect to performance metric P, based on experience E

- T: Playing checkers
- P: Percentage of games won against an arbitrary opponent
- E: Playing practice games against itself

- T: Recognizing hand-written words
- P: Percentage of words correctly classified
- E: Database of human-labeled images of handwritten words

- T: Driving on four-lane highways using vision sensors
- P: Average distance traveled before a human-judged error
- E: A sequence of images and steering commands recorded while observing a human driver.

- T: Categorize email messages as spam or legitimate.
- P: Percentage of email messages correctly classified.
- E: Database of emails, some with human-given labels

Autonomous Driving



Flying Robots

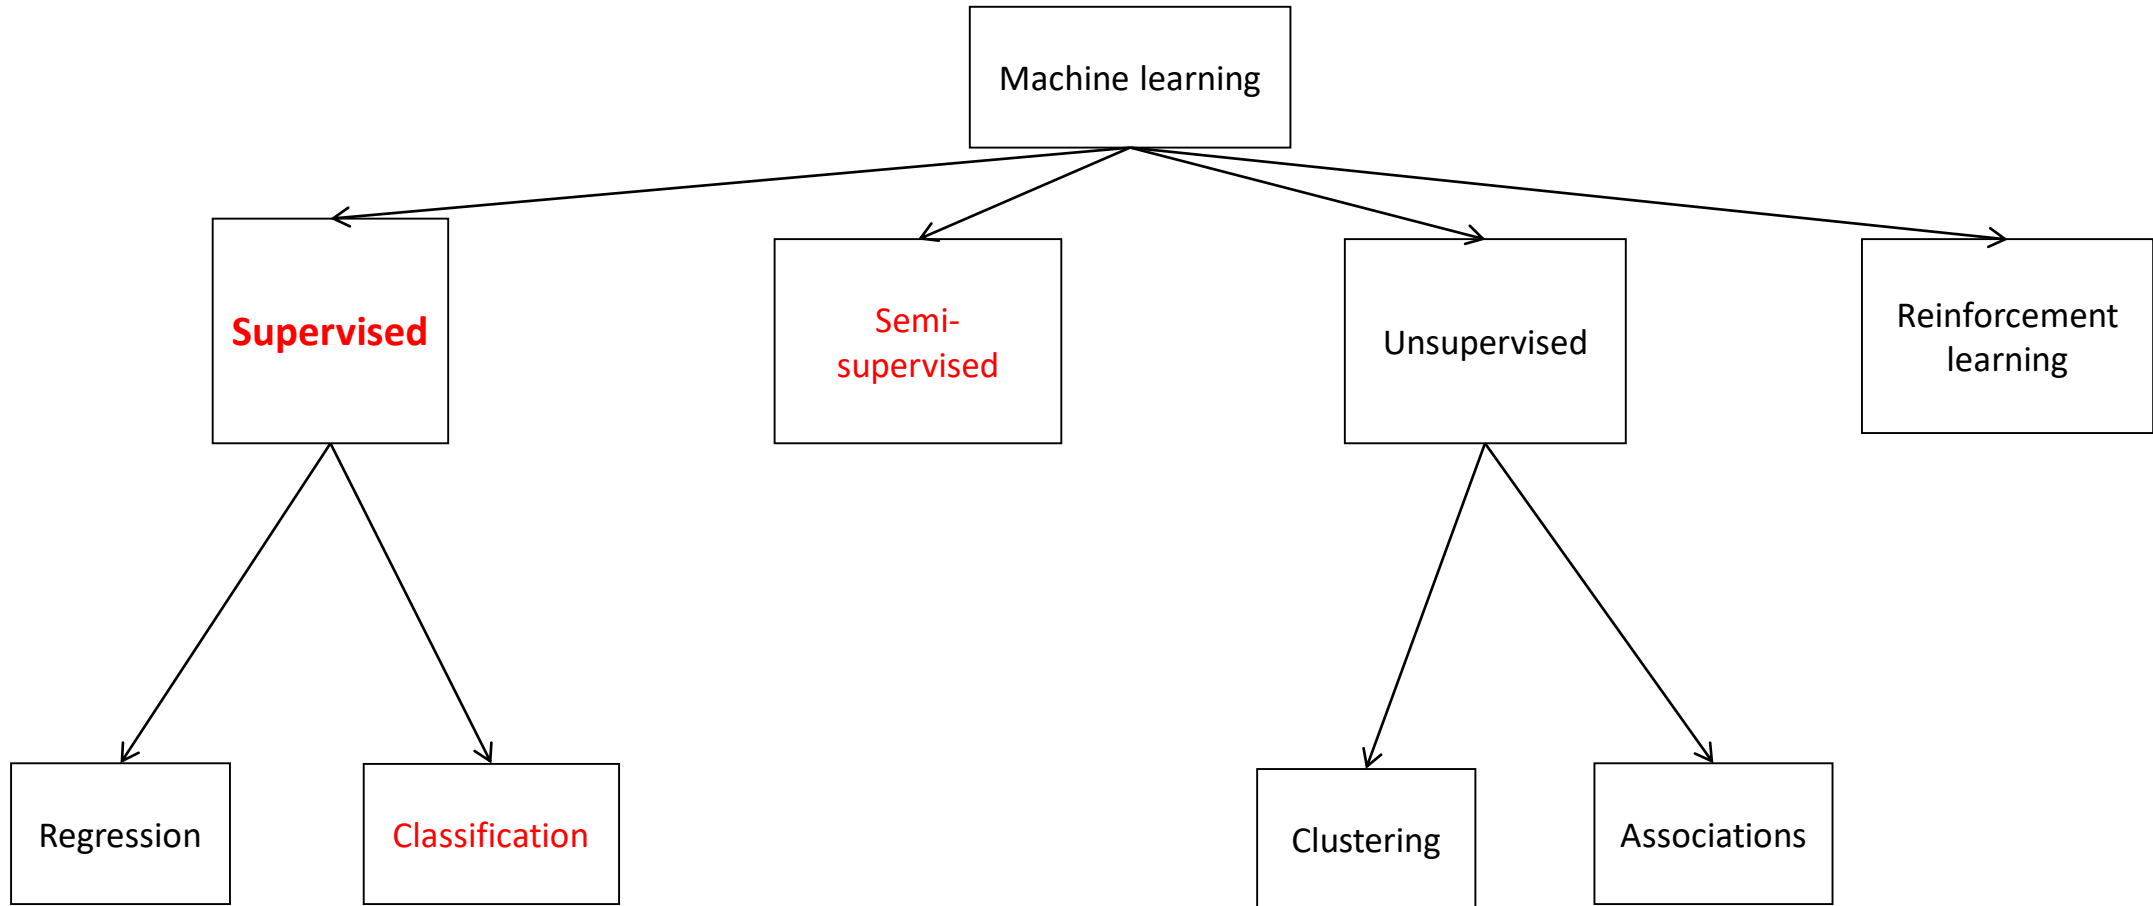


Figure: Video: <https://www.youtube.com/watch?v=YQIMGV5vtd4>

Types of Learning

- Supervised (inductive) learning
 - Given: training data + desired outputs (labels)
- Unsupervised learning
 - Given: training data (without desired outputs)
- Semi-supervised learning
 - Given: training data + a few desired outputs
- Reinforcement learning
 - Rewards from sequence of actions. An agent interacting with the world makes observations, takes actions, and is rewarded or punished; it should learn to choose actions in such a way as to obtain a lot of reward.

ML taxonomy

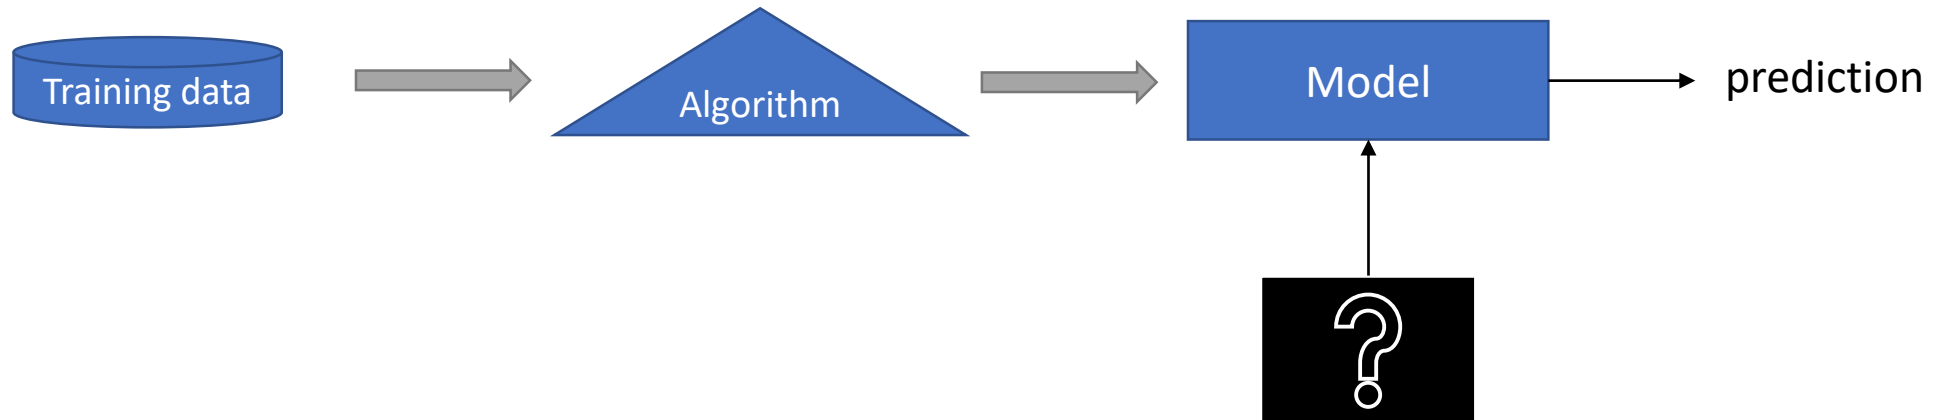


Supervised Learning

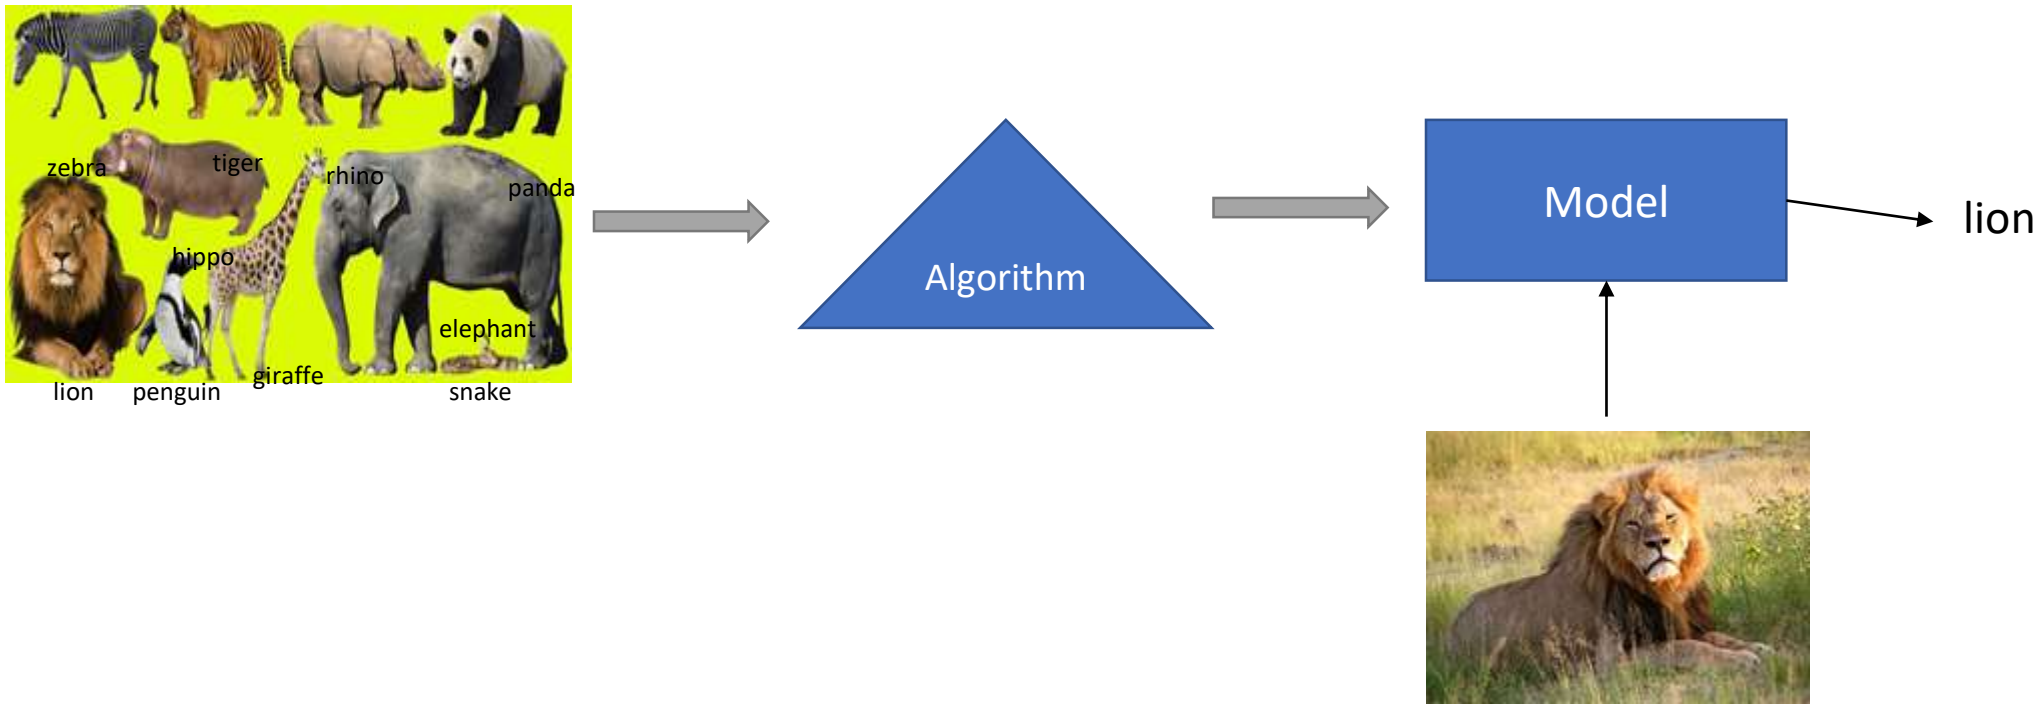
- Learning a discrete function: **Classification**
 - Boolean classification:
 - Each example is classified as true(positive) or false(negative).
 - predict categorical values, i.e., labels
- Learning a continuous function: **Regression**
 - predict numerical values

Supervised learning

- In supervised learning, the algorithms are presented with a set of classified instances from which they learn a way of classifying unseen instances. When the attribute to be predicted is numeric rather than nominal it is called regression.



Classification



Classification—A Two-Step Process

- Model construction: describing a set of predetermined classes
 - Each tuple/sample is assumed to belong to a predefined class, as determined by the class label
 - The set of tuples used for model construction is training set
 - The model is represented as **classification rules**, **decision trees**, or **mathematical formulae**
- Model usage: for classifying future or unknown objects
 - Estimate accuracy of the model
 - The known label of test sample is compared with the classified result from the model
 - **Test set is independent of training set**, otherwise over-fitting will occur
 - If the accuracy is acceptable, use the model to classify data tuples whose class labels are not known

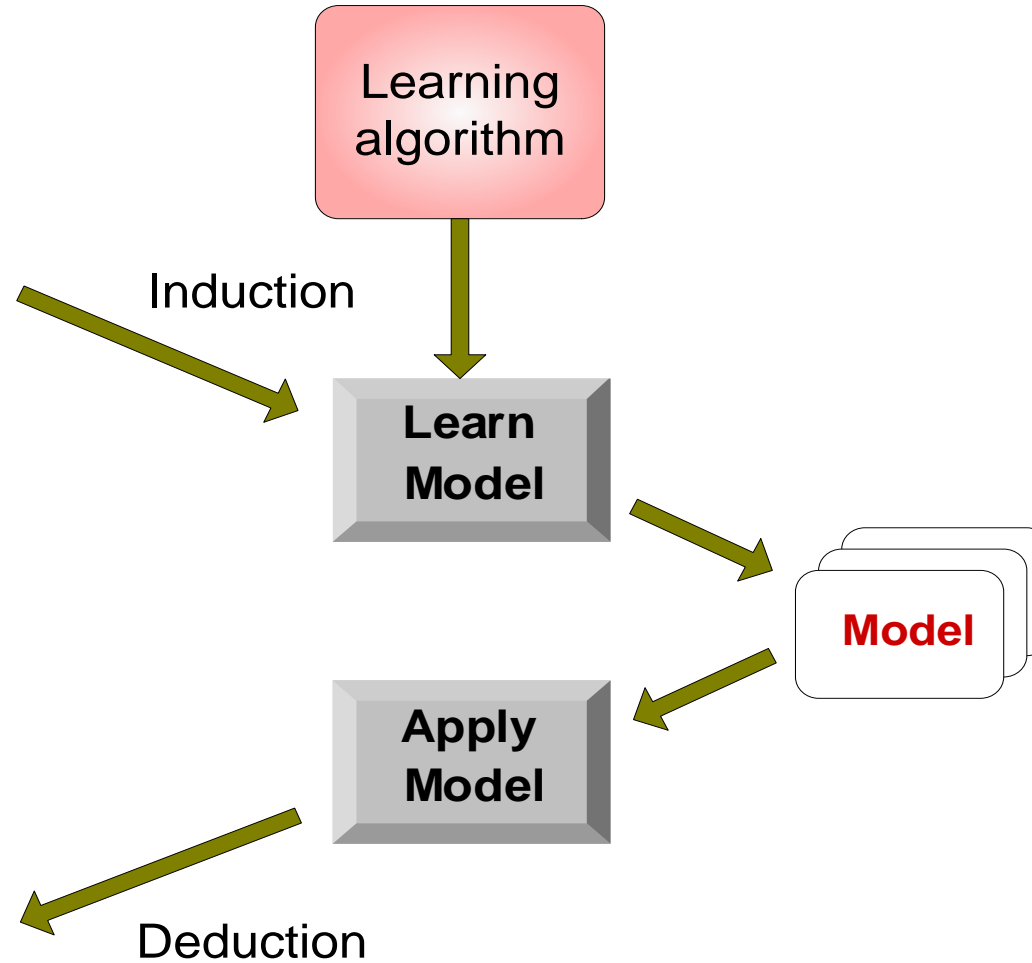
Illustrating Classification Task

<i>Tid</i>	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

<i>Tid</i>	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



Issues: Data Preparation


- Data cleaning
 - Preprocess data in order to reduce noise and handle missing values
- Relevance analysis (feature selection)
 - Remove the irrelevant or redundant attributes
- Data transformation
 - Generalize data to (higher concepts, discretization)
 - Normalize attribute values

Classification Techniques

- Decision Tree based Methods
- Rule-based Methods
- Naïve Bayes and Bayesian Belief Networks
- Neural Networks
- Support Vector Machines
- and more...

Classification

- Example: Credit scoring
- Differentiating between **low-risk** and **high-risk** customers from their *income* and *savings*



Discriminant: IF $income > \theta_1$ AND $savings > \theta_2$
THEN **low-risk** ELSE **high-risk**

Classification: Applications

- Aka Pattern recognition
- Face recognition: Pose, lighting, occlusion (glasses, beard), make-up, hair style
- Character recognition: Different handwriting styles.
- Speech recognition: Temporal dependency.
 - Use of a dictionary or the syntax of the language.
 - Sensor fusion: Combine multiple modalities; eg, visual (lip image) and acoustic for speech
- Medical diagnosis: From symptoms to illnesses
- ...

Face Recognition

Training examples of a person




Test images



Supervised Learning: Regression

- Given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function $f(x)$ to predict y given x
 - y is real-valued == regression



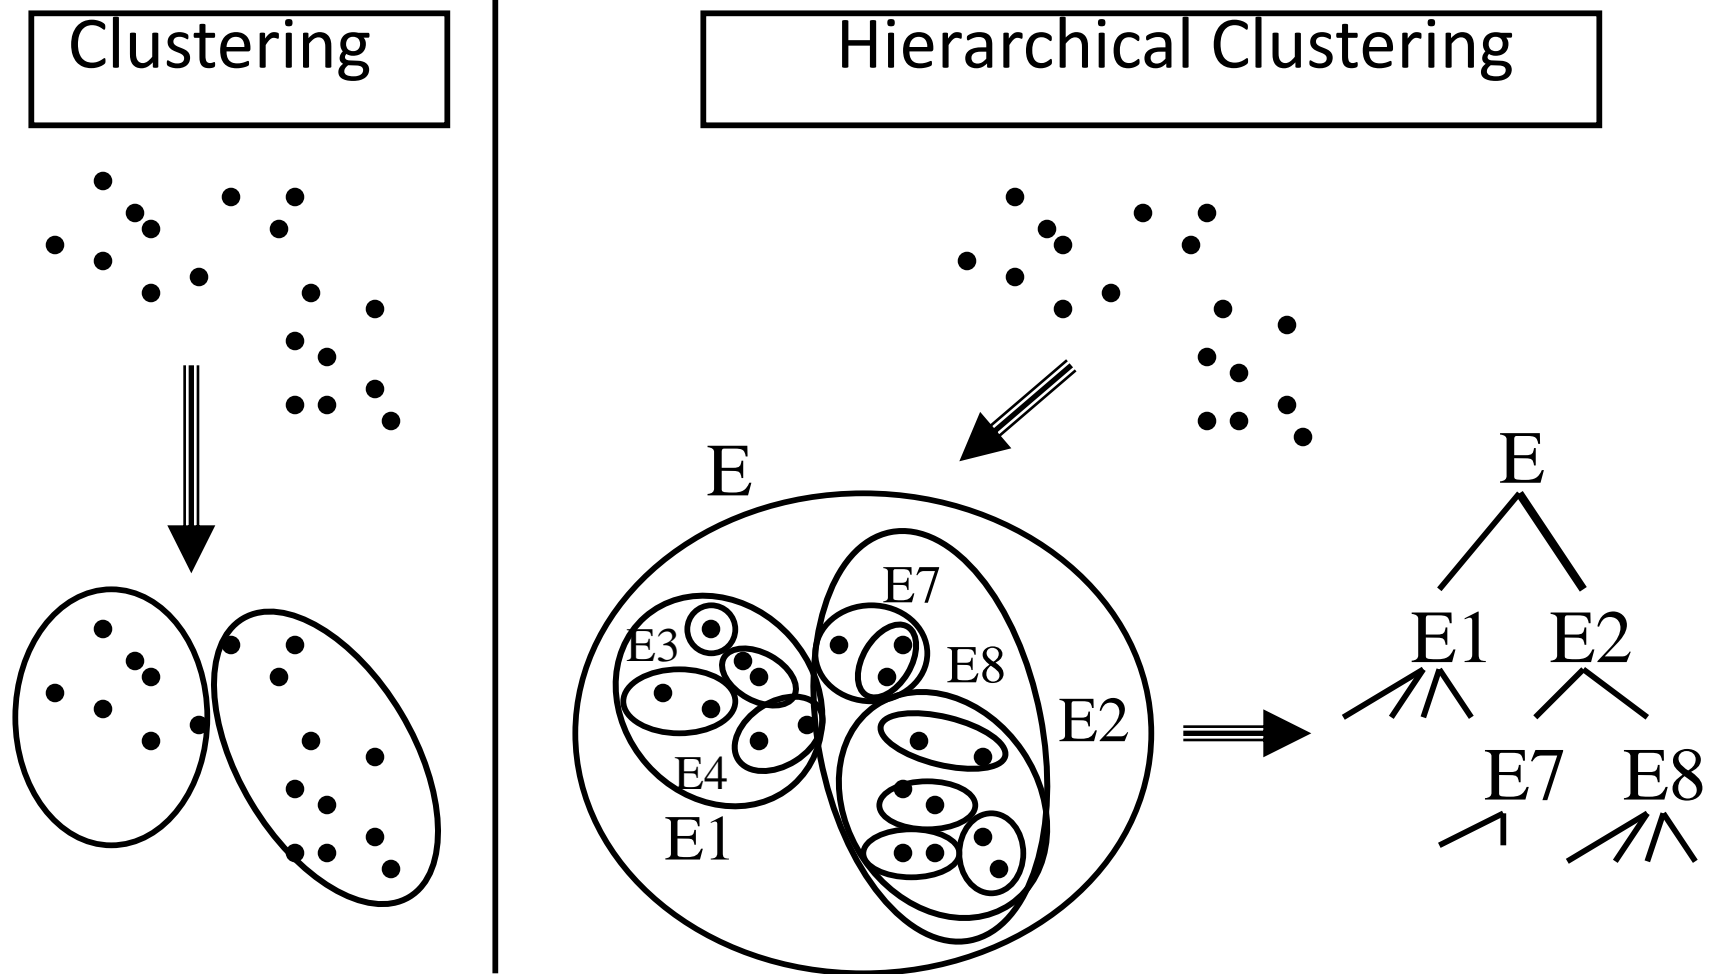
Supervised Learning: Uses

- Prediction of future cases: Use the rule to predict the output for future inputs
- Knowledge extraction: The rule is easy to understand
- Compression: The rule is simpler than the data it explains
- Outlier detection: Exceptions that are not covered by the rule, e.g., fraud

Unsupervised Learning: Overview

- So far, in all the learning techniques we considered, a training example consisted of a set of attributes (or features) and either a class (in the case of classification) or a real number (in the case of regression) attached to it.
- Unsupervised Learning takes as training examples the set of attributes/features alone.
- The purpose of unsupervised learning is to attempt to find natural partitions in the training set.
- Two general strategies for Unsupervised learning include: ***Clustering*** and ***Hierarchical Clustering***.

Clustering and Hierarchical Clustering




Other Unsupervised Methods:

- There are a lot of other Unsupervised Learning Methods.
- **Examples:**
 - k-means
 - The EM Algorithm
 - Competitive Learning
 - Kohonen's Neural Networks: Self-Organizing Maps
 - Principal Component Analysis, Autoassociation

Supervised Learning vs Unsupervised Learning

$x \rightarrow y$



x



x




y




y

Supervised Learning vs Unsupervised Learning

$x \rightarrow y$

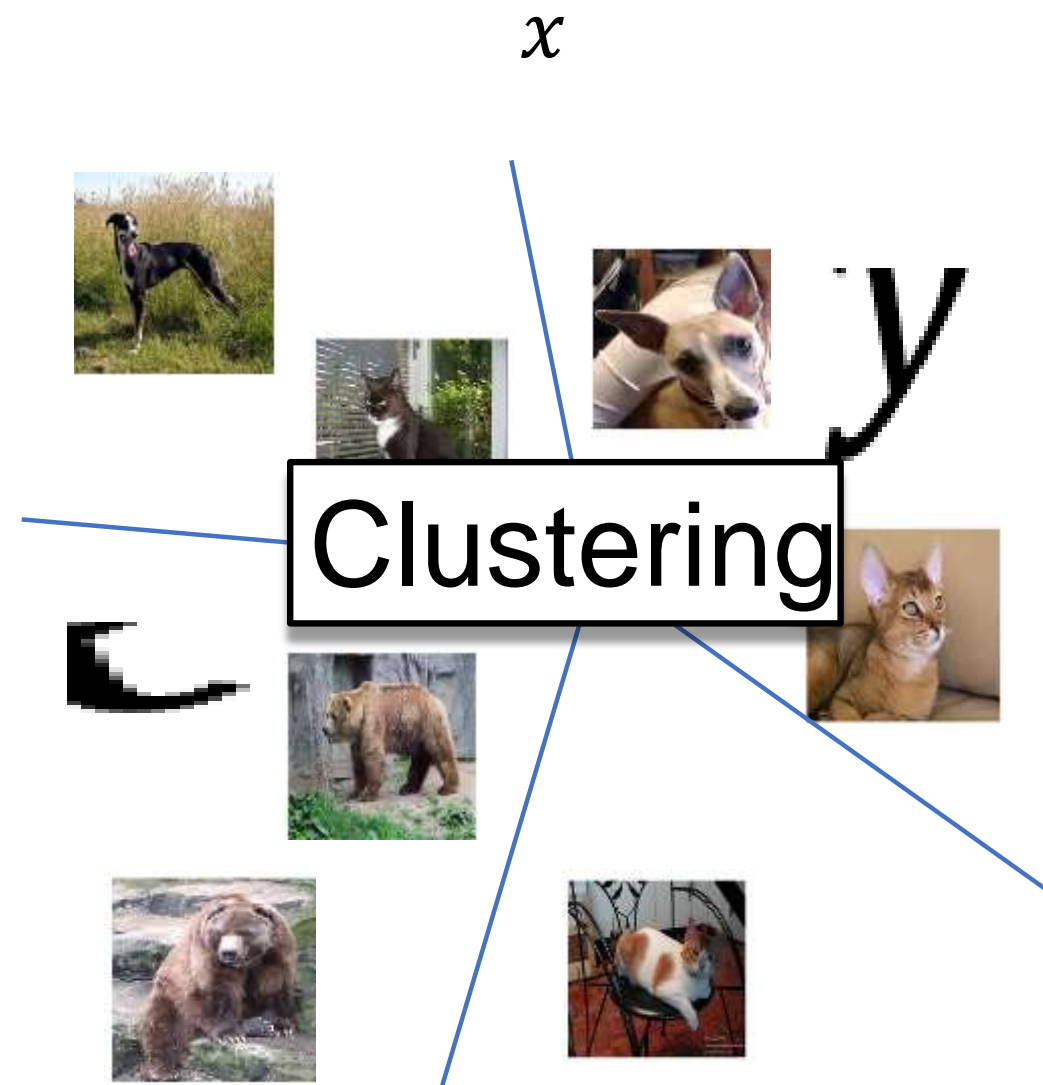
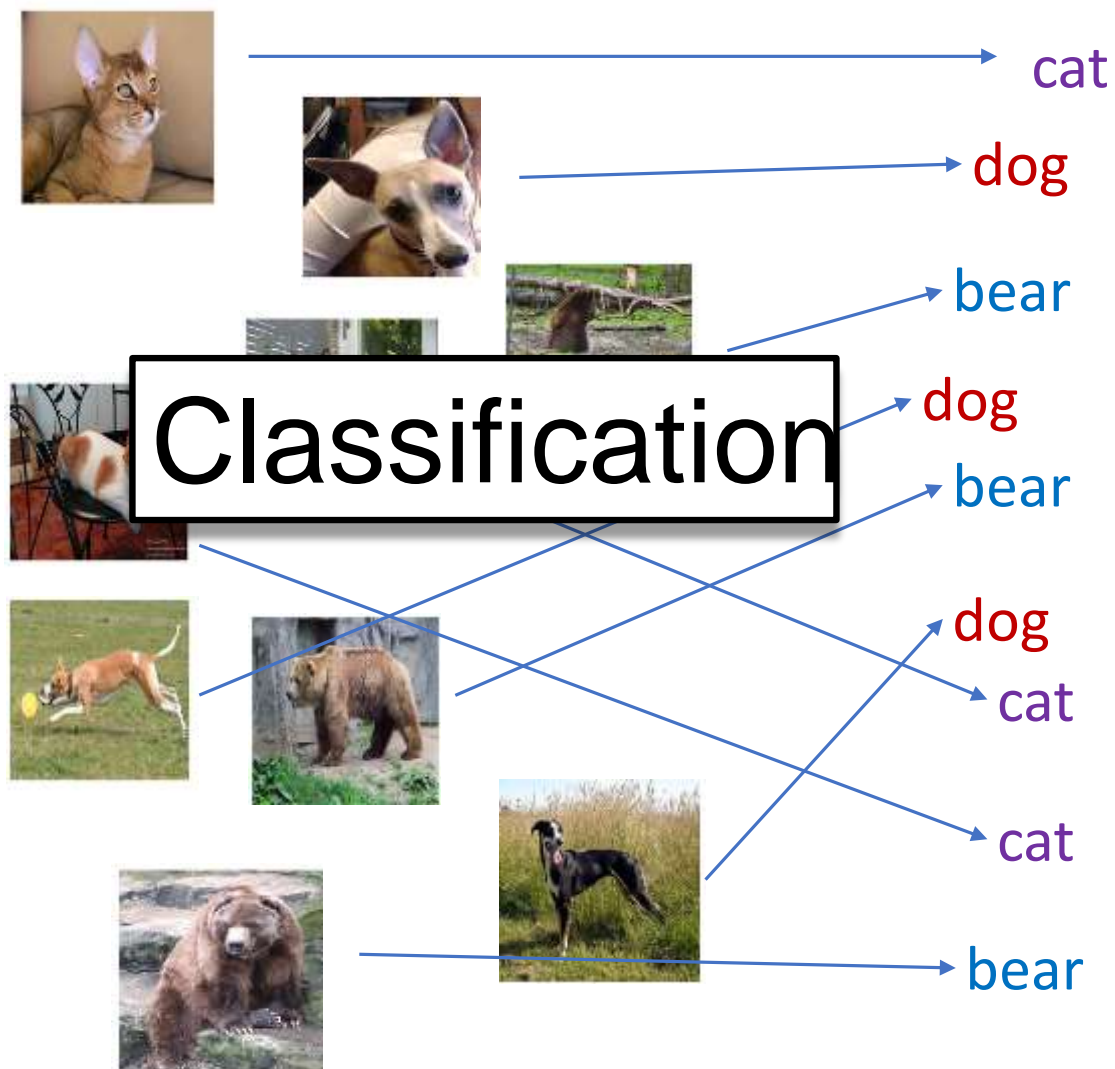


x



Supervised Learning vs Unsupervised Learning

$$x \rightarrow y$$



Reinforcement Learning

- Given a sequence of states and actions with (delayed) rewards, output a policy
 - Policy is a mapping from states to actions that tells you what to do in a given state
- Examples:
 - Credit assignment problem
 - Game playing
 - Robot in a maze
 - Balance a pole on your hand

Reinforcement Learning


Agent and environment interact at discrete time steps : $t = 0, 1, 2, \dots, K$

Agent observes state at step t : $s_t \in S$

produces action at step t : $a_t \in A(s_t)$


gets resulting reward : $r_{t+1} \in \mathcal{R}$

and resulting next state : s_{t+1}



Designing a Learning System

- Choose the training experience
- Choose exactly what is to be learned
 - i.e. the **target function**
- Choose how to represent the target function
- Choose a learning algorithm to infer the target function from the experience



Case Study

- What grade will I get in this course?
- Data: entry survey and marks from this and previous years
- Process the data
 - Split into training set; and test set
 - Determine representation of input;
 - Determine the representation of the output;
- Choose form of model: linear regression
- Decide how to evaluate the system's performance: objective function
- Set model parameters to optimize performance
- Evaluate on test set: generalization

ML in a Nutshell

- Tens of thousands of machine learning algorithms
 - Hundreds new every year
- Every ML algorithm has three components:
 - **Representation**
 - **Optimization**
 - **Evaluation**

Various Function Representations

- • Numerical functions
 - Linear regression
 - Neural networks
 - Support vector machines
- Symbolic functions
 - Decision trees
 - Rules in propositional logic
 - Rules in first-order predicate logic
- Instance-based functions
 - Nearest-neighbor
 - Case-based
- Probabilistic Graphical Models
 - Naïve Bayes
 - Bayesian networks
 - Hidden-Markov Models (HMMs)
 - Probabilistic Context Free Grammars (PCFGs)
 - Markov networks

Various Search/Optimization Algorithms

- Gradient descent
 - Perceptron
 - Backpropagation
- Dynamic Programming
 - HMM Learning
 - PCFG Learning
- Divide and Conquer
 - Decision tree induction
 - Rule learning
- Evolutionary Computation
 - Genetic Algorithms (GAs)
 - Genetic Programming (GP)
 - Neuro-evolution

Evaluation

- Accuracy
- Precision and recall
- Squared error
- Likelihood
- Posterior probability
- Cost / Utility
- Margin
- Entropy
- K-L divergence
- etc.

ML in Practice



- Understand domain, prior knowledge, and goals
- Data integration, selection, cleaning, pre-processing, etc.
- Learn models
- Interpret results
- Consolidate and deploy discovered knowledge

Lessons Learned about Learning

- Learning can be viewed as using direct or indirect experience to approximate a chosen target function.
- Function approximation can be viewed as a search through a space of hypotheses (representations of functions) for one that best fits a set of training data.
- Different learning methods assume different hypothesis spaces (representation languages) and/or employ different search techniques